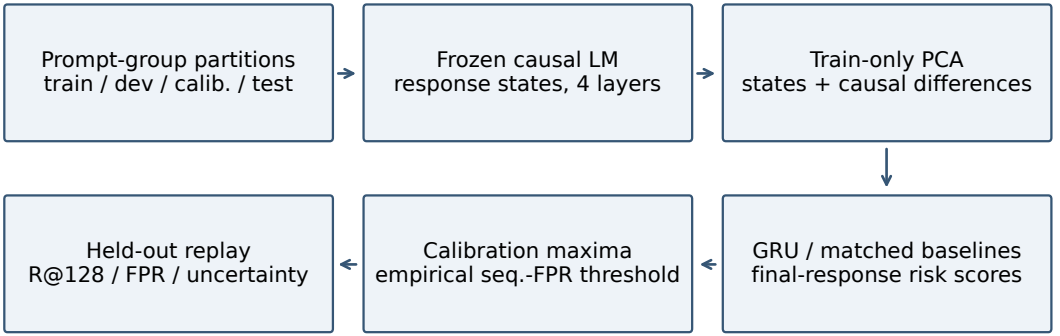
Prompt-group partitions
train / dev / calib. / test
Frozen causal LM
response states, 4 layers
Train-only PCA
states + causal differences
GRU / matched baselines
final-response risk scores
Calibration maxima
empirical seq.-FPR threshold
Held-out replay
R@128 / FPR / uncertainty